


TECHNICAL ARCHITECTURE PLAN

Hermes Agent Android App

A comprehensive technical plan for building a full AI agent platform optimized for Samsung Galaxy S24 Ultra, featuring multi-agent orchestration, on-device LLM inference, plugin ecosystem, and hardware-accelerated intelligence.

Target Device: Samsung Galaxy S24 Ultra
Platform: Android 14+ (Kotlin Native)
Version: 1.0 Planning Document



JUNE 2026

Table of Contents

1. Executive Summary	3
2. Product Vision & Goals	3
2.1 Vision Statement	3
2.2 Target Users	4
2.3 Key Objectives	4
3. System Architecture	4
3.1 High-Level Architecture	4
3.2 Agent Orchestration Flow	5
3.3 Plugin System Architecture	6
4. Technology Stack	6
4.1 Core Technology Recommendation	6
4.2 Key Technology Components	7
5. Samsung S24 Ultra Optimization	8
5.1 NPU Acceleration for On-Device Inference	8
5.2 Memory Management Strategy	8
5.3 S Pen SDK and Input Modalities	9
5.4 Battery Optimization	9
6. Core Features - Agent Platform	9
6.1 Multi-Agent Orchestration	9
6.2 Conversation Memory System	10
6.3 RAG Pipeline	10
6.4 Feature Priority Matrix	11
7. Development Roadmap	11

7.1 Phase 1: Foundation (Weeks 1-6)	11
7.2 Phase 2: Core Agent (Weeks 7-14)	12
7.3 Phase 3: Platform Features (Weeks 15-20)	12
7.4 Phase 4: Polish and Launch (Weeks 21-24)	12
8. Risk Assessment & Mitigation	13
8.1 Technical Risks	13
8.2 Product and Market Risks	14

1. Executive Summary

Hermes Agent is envisioned as a full-featured AI agent platform designed from the ground up for the Samsung Galaxy S24 Ultra. Unlike conventional chatbot applications that provide simple question-and-answer interactions, Hermes Agent functions as an intelligent personal companion capable of orchestrating multiple specialized agents, executing complex multi-step tasks, maintaining persistent memory across conversations, and leveraging both on-device and cloud-based large language models (LLMs) to deliver a seamless and privacy-respecting user experience.

The platform architecture centers on a multi-agent orchestration engine that delegates user requests to specialized sub-agents, each optimized for specific domains such as productivity, information retrieval, communication, and device control. A plugin system enables third-party developers and power users to extend agent capabilities, creating an open ecosystem around core intelligence features. The Retrieval-Augmented Generation (RAG) pipeline combines on-device vector search with optional cloud retrieval, ensuring that the agent can reference personal documents, notes, and contextual knowledge without compromising data privacy.

A distinguishing feature of Hermes Agent is its deep integration with the S24 Ultra hardware capabilities. The Snapdragon 8 Gen 3 Neural Processing Unit (NPU) enables accelerated on-device LLM inference, while 12 GB of RAM provides the memory headroom necessary for running language models locally alongside the Android operating system. Samsung Knox security features ensure that all agent operations, including local model inference and data storage, operate within a hardened sandbox environment. S Pen SDK integration allows handwritten input to be processed directly by the agent, opening new interaction modalities beyond traditional keyboard and voice.

2. Product Vision & Goals

2.1 Vision Statement

The vision for Hermes Agent is to transform the smartphone from a passive tool that requires explicit user instructions into an active, context-aware intelligence partner. Rather than forcing users to navigate multiple applications and manually coordinate information across services, Hermes Agent understands intent, manages context across conversations, and proactively suggests or executes actions on the user's behalf. The S24 Ultra, with its powerful NPU and ample memory, represents the first generation of consumer devices where meaningful on-device AI agent capabilities become practically feasible.

The name "Hermes" draws from the Greek messenger god, symbolizing the app's role as an intelligent intermediary between the user and the vast ecosystem of digital services, device capabilities, and knowledge sources. Hermes does not replace human decision-making; it amplifies human capability by handling the cognitive overhead of task execution, information synthesis, and cross-application coordination that consumes a significant portion of daily digital interaction time.

2.2 Target Users

The primary target audience for Hermes Agent consists of power users and technology professionals who rely heavily on their smartphones for productivity, communication, and information management. These users typically juggle multiple communication channels, manage complex schedules, work with large volumes of documents and notes, and frequently need to retrieve, synthesize, and act upon information from diverse sources. A secondary audience includes developers and AI enthusiasts who are interested in extending agent capabilities through the plugin system and experimenting with on-device AI workflows.

2.3 Key Objectives

Objective	Category
Deliver a fully functional multi-agent orchestration platform on-device	v1.0 Launch
Achieve sub-3-second response time for on-device inference tasks	Performance
Enable privacy-first architecture where sensitive data never leaves the device	Security
Support a minimum of 20 first-party tool integrations at launch	Ecosystem
Maintain battery impact below 8% per day of active agent usage	Efficiency
Establish a plugin marketplace framework for third-party extensions	Platform

Table 1: Key objectives and success categories for Hermes Agent v1.0

3. System Architecture

3.1 High-Level Architecture

The Hermes Agent architecture follows a layered design pattern that cleanly separates user interface concerns from agent intelligence and data management responsibilities. At the top layer, the presentation layer built with Jetpack Compose provides a responsive and modern Android user interface supporting chat interactions, voice input and output, S Pen integration, and system notification-based agent alerts. This layer communicates exclusively with the Agent Orchestration layer through a well-defined Kotlin coroutine-based API, ensuring that UI updates remain decoupled from potentially long-running inference operations.

The Agent Orchestration layer serves as the central nervous system of the application. It receives user requests, classifies intent, routes tasks to appropriate specialized agents, manages conversation state, and coordinates tool execution. The Multi-Agent Manager maintains a registry of available agents, each with defined capabilities, resource requirements, and priority levels. A conversation memory subsystem separates short-term working memory (current session context, recent tool results) from long-term

semantic memory (user preferences, learned patterns, historical summaries stored in a local vectordatabase).

The Intelligence and Data layer provides the foundational AI capabilities. On-device LLM inference runs through MLC-LLM or a llama.cpp-based runtime that leverages the Snapdragon 8 Gen 3 NPU via the Qualcomm AI Engine Direct SDK. For tasks requiring greater model capacity, a cloud LLM fallback connects to remote API endpoints (OpenAI, Anthropic, or self-hosted models) through an encrypted channel. The RAG pipeline indexes user documents, notes, and conversation history into a local vector store (SQLite-VSS or ChromaDB embedded) and retrieves relevant context at query time to augment LLM prompts.




Figure 1: Hermes Agent system architecture showing the layered design from user interface through agent orchestration to intelligence and hardware layers

3.2 Agent Orchestration Flow

When a user submits a query or command, the orchestration pipeline follows a deterministic yet flexible flow. The Agent Router first performs intent classification using a lightweight on-device model to categorize the request type (conversational, task-execution, information retrieval, device control, or multi-step planning). Based on this classification, the router selects one or more agents from the registry and constructs an execution plan. For simple queries, a single agent handles the request end-to-end. For complex, multi-domain tasks, the router decomposes the request into sub-tasks, assigns each to the most

appropriate agent, and defines the data flow and dependencies between sub-task results.

The Tool Registry maintains a catalog of available tools organized by category: device APIs (contacts, calendar, camera, file system, notifications), communication tools (SMS, email, messaging apps), information tools (web search, document retrieval, knowledge base queries), and productivity tools (task management, scheduling, reminders). Each tool is registered with a JSON schema describing its input parameters, output format, required permissions, and execution constraints. Agents invoke tools through a standardized function-calling protocol compatible with OpenAI-style function definitions, enabling seamless integration with both local and cloud LLM providers.

3.3 Plugin System Architecture

The plugin system follows a sandboxed extension model where each plugin operates within its own isolated execution environment. Plugins are packaged as Android APK modules with a defined interface contract that exposes capabilities to the agent orchestration layer. The interface contract specifies four required components: a capability descriptor (what the plugin can do), a tool schema (input/output definitions for function calling), a permission manifest (device resources the plugin requires), and a lifecycle handler (initialization, suspension, and cleanup callbacks). Plugins communicate with the core application through a gRPC-based inter-process communication channel, ensuring that a misbehaving plugin cannot crash or compromise the main agent process.

The plugin marketplace framework provides a discovery, installation, and update mechanism similar to modern app store patterns. Plugins are signed with developer certificates and verified against a schema compliance check before installation. Runtime resource monitoring enforces CPU, memory, and network usage limits per plugin, preventing any single extension from degrading overall system performance. A permissions review screen presents users with a clear summary of what each plugin can access before installation, maintaining the privacy-first architecture principle.

4. Technology Stack

4.1 Core Technology Recommendation

After evaluating the three primary mobile development approaches, the recommended technology stack for Hermes Agent is Kotlin native development with Jetpack Compose for the UI layer. This recommendation is driven by several critical factors specific to the project requirements. First, deep Samsung hardware integration (NPU acceleration, S Pen SDK, Knox security APIs) requires access to native Android APIs and Samsung-specific SDKs that are not fully exposed through cross-platform frameworks. Second, on-device LLM inference through MLC-LLM or llama.cpp operates through JNI bindings that are most naturally integrated from Kotlin or Java code. Third, performance-critical operations such as real-time vector search, memory management, and concurrent agent execution benefit from Kotlin coroutine support and direct access to Android's low-level optimization APIs.

Framework	Key Libraries	Hardware Integration	Recommendation	Caveat
Kotlin Native	Jetpack Compose, Coroutines, Room	Excellent - full hardware API access, direct NPU/JNI integration, best performance	Recommended	Steeper Android-specific learning curve
React Native + Hermes	React Navigation, AsyncStorage	Good - NPU via native modules, Hermes engine for JS performance	Viable	NPU/LLM integration requires native bridge modules
Flutter	Dart, FFI, Isolates	Moderate - NPU via platform channels, FFI for C++ libs	Not Recommended	Platform channel overhead for real-time AI workloads

Table 2: Technology stack comparison for Hermes Agent development

4.2 Key Technology Components

Component	Technology	Purpose
UI Framework	Jetpack Compose	Declarative UI, Material 3 design, reactive state management
Architecture	MVVM + Clean Architecture	Separation of presentation, domain, and data layers
Dependency Injection	Hilt (Dagger)	Compile-time DI, scoped instances for agents and tools
Local Database	Room + SQLite-VSS	Structured storage with embedded vector search for RAG
On-Device LLM	MLC-LLM / llama.cpp	Quantized model inference via NPU acceleration
Cloud LLM	OpenAI / Anthropic SDK	Fallback inference for complex reasoning tasks
Agent Framework	Custom (inspired by LangChain)	Kotlin-native agent orchestration and tool calling
Vector Store	ChromaDB Embedded	Document embedding storage and similarity search
IPC	gRPC	Inter-process communication for plugin sandboxing
Networking	OkHttp + Retrofit	REST and streaming API communication
Serialization	Kotlinx Serialization	JSON, Protocol Buffers, and custom format support
Security	Samsung Knox + Android Keystore	Hardware-backed key storage and sandbox isolation

5. Samsung S24 Ultra Optimization

5.1 NPU Acceleration for On-Device Inference

The Snapdragon 8 Gen 3 processor in the Galaxy S24 Ultra features a Qualcomm Hexagon NPU capable of delivering up to 45 TOPS (Trillions of Operations Per Second) of AI performance. This represents a significant leap over previous generations and makes on-device inference of quantized large language models practically feasible for the first time in a consumer smartphone. Hermes Agent leverages the Qualcomm AI Engine Direct SDK to delegate model inference operations to the NPU, achieving inference speeds 3-5x faster than CPU-only execution while consuming significantly less power. The MLC-LLM runtime provides pre-compiled model artifacts optimized for the Adreno GPU and Hexagon NPU, using 4-bit quantized variants of models like Llama-3-8B or Phi-3-mini that fit within the S24 Ultra's 12 GB RAM allocation while maintaining acceptable quality for agent tasks.

The inference pipeline implements a dynamic routing strategy that evaluates each request against a complexity threshold. Simple queries (factoid questions, device control commands, template-based responses) are handled entirely on-device with zero latency from network calls. Complex reasoning tasks (multi-step planning, nuanced analysis, creative generation) are routed to cloud LLM endpoints when the user has opted in to cloud-assisted mode. A hybrid approach splits the workload: the on-device model handles initial intent classification and tool selection, while the cloud model provides the final generation for demanding tasks. This strategy reduces cloud API costs by an estimated 60-70% compared to routing all requests to remote endpoints.

5.2 Memory Management Strategy

With 12 GB of LPDDR5X RAM, the S24 Ultra provides substantial memory headroom, but Android's memory management system imposes strict limits on background process memory usage to maintain system responsiveness. Hermes Agent implements a tiered memory management strategy that ensures the on-device LLM, vector database, and agent runtime coexist peacefully with the operating system and other applications. The LLM runtime occupies approximately 2.5-3.5 GB of RAM depending on the loaded model size, while the vector database and conversation memory subsystem consume an additional 500 MB to 1 GB during active sessions.

To prevent the system from killing the agent process during memory pressure events, Hermes implements a proactive memory shedding strategy. When the system reports available memory below a 2 GB threshold, the agent gracefully offloads the on-device LLM weights to flash storage, retains only the vector database index in memory for RAG queries, and switches all inference to cloud mode until memory pressure subsides. This approach ensures that the agent remains functional even under heavy multitasking conditions, trading inference speed for resilience. The memory state machine is fully observable through a diagnostic panel in the app settings, allowing power users to monitor real-time

memory allocation and configure shedding thresholds.

5.3 S Pen SDK and Input Modalities

The S Pen on the S24 Ultra presents a unique input modality that Hermes Agent is designed to leverage extensively. Through the Samsung S Pen SDK, the application receives real-time stroke data including pressure sensitivity, tilt angle, and velocity, enabling the agent to interpret handwritten notes, diagrams, and annotations as natural input. A lightweight on-device handwriting recognition model converts S Pen input to digital text before passing it through the standard agent pipeline, while raw stroke data is preserved for tasks that benefit from spatial understanding, such as diagram interpretation or mathematical equation solving.

5.4 Battery Optimization

Always-on agent capabilities require careful battery management to avoid unacceptable drain. Hermes employs several optimization strategies: NPU inference is prioritized over GPU and CPU because it delivers the best performance-per-watt ratio for neural network operations. Background agent monitoring (notification processing, scheduled task execution) uses Android's WorkManager API with constrained constraints that batch operations during device charging or maintenance windows. The on-device LLM model is loaded lazily and unloaded after a configurable idle period (default: 5 minutes of no agent interaction), reducing standby power consumption to near-zero. Field testing targets indicate that with mixed on-device and cloud inference usage patterns, the expected daily battery impact is 5-8%, well within acceptable limits for a power-user application.

6. Core Features - Agent Platform

6.1 Multi-Agent Orchestration

The multi-agent orchestration system is the intellectual core of Hermes Agent. Rather than relying on a single monolithic model to handle all user requests, the system decomposes complex tasks across a team of specialized agents, each fine-tuned or prompted for a specific domain. The default agent roster includes a Conversational Agent (natural dialogue and small talk), a Productivity Agent (calendar management, task creation, email drafting), a Research Agent (web search, document analysis, summarization), a Device Control Agent (system settings, app launching, notification management), and a Creative Agent (writing assistance, brainstorming, content generation). Each agent maintains its own context window, tool access permissions, and response formatting preferences, allowing them to operate semi-independently while the orchestrator manages inter-agent communication and result aggregation.

The orchestration algorithm follows a plan-then-execute pattern inspired by recent advances in agentic AI architectures. When a user request is received, the orchestrator first generates a high-level execution plan that identifies which agents are needed, what tools each agent should invoke, and how intermediate results should flow between agents. This plan is presented to the user as a brief, editable outline before

execution begins, giving users transparency and control over the agent's intended actions. During execution, the orchestrator monitors each agent's output, handles errors and retries, and can dynamically adjust the plan if intermediate results suggest a different approach would be more effective.

6.2 Conversation Memory System

The memory system implements a dual-store architecture inspired by cognitive science models of human memory. Short-term episodic memory captures the current conversation context, including recent user statements, agent responses, tool invocations, and their results. This memory is maintained as a sliding window of token-limited conversation turns (typically the last 20-30 exchanges) that form the immediate context for LLM inference. Long-term semantic memory persists across sessions and captures user preferences, frequently referenced information, learned patterns, and summaries of past conversations. Semantic memories are stored as embeddings in the local vector database, indexed by relevance and recency, and retrieved during inference to provide personalized context even when the current conversation does not explicitly reference past information.

A memory consolidation process runs during device charging to compress and organize accumulated episodic memories into semantic memories. This process identifies recurring themes, extracts key facts and preferences, and generates concise summaries that preserve essential information while reducing storage requirements. Users can review, edit, and delete individual memories through a dedicated memory management interface, maintaining full control over what the agent remembers. The consolidation algorithm also prunes low-relevance or outdated memories based on a recency-weighted relevance score, preventing unbounded memory growth while preserving the most valuable information.

6.3 RAG Pipeline

The Retrieval-Augmented Generation pipeline enables Hermes Agent to ground its responses in the user's personal data and documents. The pipeline operates in three stages: ingestion, indexing, and retrieval. During ingestion, documents (PDF files, notes, web pages, email threads) are chunked into semantically meaningful segments using a recursive text splitter that respects paragraph and section boundaries. Each chunk is embedded using a small on-device embedding model (such as all-MiniLM-L6-v2 quantized) and stored in the local vector database alongside metadata including source document identifier, chunk position, creation timestamp, and access frequency. During retrieval, the user's query is embedded using the same model, and the vector database returns the top-k most similar chunks using approximate nearest neighbor search.

The RAG pipeline supports a hybrid retrieval strategy that combines vector similarity search with keyword-based BM25 retrieval to handle both semantic and exact-match queries effectively. Retrieved chunks are reranked using a cross-encoder model that evaluates the relevance of each chunk in the context of the full query, ensuring that the most useful context is provided to the LLM. For users who opt in to cloud-assisted mode, the pipeline can also query remote knowledge bases and web search results, merging on-device and cloud retrieval results into a unified context window. An intelligent context window manager ensures that retrieved content fits within the model's token limit while

maximizing information density by prioritizing diverse, non-redundant chunks.

6.4 Feature Priority Matrix

Feature	Priority	Rationale
Multi-Agent Orchestration	P0 - Must Have	Core differentiator, defines the platform
On-Device LLM Inference	P0 - Must Have	Privacy foundation, enables offline operation
Conversation Memory	P0 - Must Have	Essential for contextual, personalized interactions
RAG Pipeline	P0 - Must Have	Grounds responses in user data and documents
Plugin System Framework	P1 - Should Have	Extensibility for long-term ecosystem growth
Cloud LLM Fallback	P1 - Should Have	Handles complex tasks beyond on-device capacity
S Pen Integration	P2 - Nice to Have	Unique S24 Ultra differentiator, not blocking
Voice Input/Output	P1 - Should Have	Key interaction modality for mobile agent usage
Notification Agent	P2 - Nice to Have	Proactive agent actions from system notifications

Table 4: Feature priority matrix for Hermes Agent v1.0

7. Development Roadmap

The development roadmap is organized into four phases spanning approximately 24 weeks total, progressing from foundational infrastructure through core agent capabilities, platform features, and launch preparation. Each phase includes defined milestones, deliverables, and quality gates that must be passed before advancing to the next phase. The timeline assumes a core development team of 3-5 engineers with specialized expertise in Android development, machine learning, and systems architecture.

7.1 Phase 1: Foundation (Weeks 1-6)

Phase 1 establishes the technical foundation upon which all agent capabilities are built. This phase focuses on setting up the project architecture, implementing the core UI shell with Jetpack Compose, integrating the on-device LLM runtime (MLC-LLM) with NPU acceleration, and building the basic chat

interface. The foundation also includes the local database layer (Room + SQLite-VSS) for conversation storage and initial vector indexing, the dependency injection framework (Hilt), and the security infrastructure leveraging Samsung Knox APIs. Key deliverables include a working chat interface with on-device model inference, basic conversation persistence, and a verified NPU acceleration pipeline benchmarked against CPU-only baseline performance.

7.2 Phase 2: Core Agent (Weeks 7-14)

Phase 2 implements the intelligence core of the platform. This is the most technically demanding phase, encompassing the agent orchestration engine, multi-agent manager, tool registry and execution framework, function calling protocol, conversation memory system (short-term and long-term), and the RAG pipeline with document ingestion and retrieval. The agent orchestration engine is built incrementally, starting with a single-agent prototype and progressively adding routing logic, plan generation, and inter-agent coordination. The RAG pipeline is developed in parallel, beginning with the embedding and indexing pipeline and progressing through retrieval, reranking, and context window management. Testing during this phase focuses heavily on agent reasoning quality, tool execution reliability, and memory retrieval accuracy.

7.3 Phase 3: Platform Features (Weeks 15-20)

Phase 3 expands the core agent capabilities into a full platform. The plugin system framework is implemented, including the sandbox architecture, gRPC communication layer, plugin lifecycle management, and the first set of first-party plugins (calendar, contacts, file manager, web browser). The cloud LLM fallback is integrated with a unified inference router that dynamically selects between on-device and cloud providers based on task complexity, user preferences, and network availability. Voice input/output is added using Android's SpeechRecognizer API for input and a lightweight TTS engine for spoken responses. Samsung-specific integrations including S Pen input processing and Galaxy AI interoperability points are completed during this phase.

7.4 Phase 4: Polish and Launch (Weeks 21-24)

The final phase focuses on performance optimization, security hardening, user experience polish, and launch preparation. Performance optimization targets include reducing cold-start latency, optimizing memory usage patterns, implementing aggressive battery saving strategies, and stress-testing the agent orchestration system under concurrent multi-agent workloads. Security hardening includes penetration testing of the plugin sandbox, encryption of all stored data at rest, and verification of Samsung Knox integration. UX polish addresses accessibility, localization for key markets, onboarding flow optimization, and visual design refinements. The phase concludes with beta testing, bug fixes, performance regression testing, and preparation of App Store listing materials including screenshots, descriptions, and privacy policy documentation.

Phase	Timeline	Key Activities	Milestone Deliverable
Phase 1: Foundation	Weeks 1-6	Project setup, UI shell, on-device LLM, NPU integration, basic chat	Working chat with local inference
Phase 2: Core Agent	Weeks 7-14	Agent orchestration, memory system, RAG pipeline, tool registry	Multi-agent system with RAG
Phase 3: Platform	Weeks 15-20	Plugin framework, cloud LLM, voice I/O, S Pen, device integrations	Full platform with extensions
Phase 4: Launch	Weeks 21-24	Performance tuning, security audit, UX polish, beta testing	Production-ready v1.0 release

Table 5: Development roadmap phases, timelines, and milestones

8. Risk Assessment & Mitigation

Building a sophisticated on-device AI agent platform for a specific hardware target introduces several categories of risk that must be proactively identified and mitigated. The following assessment covers technical, product, and security risks, each evaluated by likelihood and potential impact, with specific mitigation strategies designed to reduce residual risk to acceptable levels. The risk landscape is particularly dynamic in the mobile AI space, where hardware capabilities, model architectures, and user expectations are all evolving rapidly.

8.1 Technical Risks

The most significant technical risk is on-device LLM performance falling short of user expectations. While the Snapdragon 8 Gen 3 NPU is powerful by mobile standards, it remains orders of magnitude less capable than server-class GPUs. Quantized 4-bit models in the 7-8 billion parameter range produce noticeably lower quality outputs compared to their full-precision cloud counterparts, particularly for complex reasoning, creative writing, and nuanced multi-turn conversations. Thermal throttling during sustained inference sessions further degrades performance. The mitigation strategy centers on intelligent request routing that minimizes on-device inference for tasks where quality is critical, aggressive model optimization (quantization, pruning, speculative decoding), and setting clear user expectations through UI indicators that distinguish on-device from cloud-generated responses.

Memory management represents another critical technical risk. Loading a 4-bit quantized 8B parameter model consumes approximately 3.5 GB of RAM, leaving limited headroom for the vector database, agent runtime, and Android system processes. If the system experiences memory pressure from other applications, Android may kill the agent's background process, losing in-memory state and interrupting ongoing operations. The tiered memory shedding strategy described in Section 5.2 addresses this risk, but it introduces complexity and the potential for degraded user experience during memory-constrained periods. Continuous profiling and optimization throughout development is essential to ensure that memory usage patterns are well-understood and controlled.

8.2 Product and Market Risks

Samsung's own Galaxy AI features represent a significant competitive risk. Galaxy AI already provides on-device translation, summarization, and image editing capabilities, and Samsung is likely to expand these features significantly with future One UI updates. Hermes Agent must differentiate by offering a fundamentally more capable and extensible agent experience rather than competing on individual feature parity. The multi-agent architecture, plugin ecosystem, and persistent memory system provide this differentiation, but communicating the value of these architectural advantages to mainstream users requires careful product positioning and UX design that makes agent capabilities tangible and immediately valuable.

User adoption risk is heightened by the complexity of agent interactions compared to traditional app interfaces. Users accustomed to tapping icons and following predetermined workflows may struggle to adapt to a conversational interface where the agent's capabilities are not immediately obvious. Mitigation strategies include a carefully designed onboarding experience that progressively introduces agent capabilities through guided scenarios, a suggestion system that proactively demonstrates what the agent can do, and a visual task history that helps users understand the agent's reasoning and actions after the fact.

Risk	Likelihood	Impact	Mitigation Strategy
On-device LLM quality gap	High	High	Hybrid routing, model optimization, clear user expectations
Thermal throttling	Medium	Medium	NPU prioritization, batched inference, thermal monitoring
Memory pressure kills	Medium	High	Tiered memory shedding, process persistence, state recovery
Galaxy AI competition	High	Medium	Multi-agent differentiation, plugin ecosystem, deep integration
Plugin security breach	Low	High	Sandboxed IPC, permission model, code signing, resource limits
User adoption friction	Medium	Medium	Guided onboarding, suggestion system, task transparency
Model update management	Medium	Low	Background downloads, versioned model storage, rollback support

Table 6: Risk assessment matrix with likelihood, impact, and mitigation strategies